

SentinelAPI

Autonomous Differential API Security Testing

While 200 teams build generic AI wrappers or test basic SQLi/XSS, SentinelAPI solves the single most dangerous vulnerability in cloud software: **Broken Object Level Authorization (BOLA/IDOR)**.

WHAT TRACK C DEMANDS

OWASP API #1 Focus

Automated detection of authorization logic flaws across multi-tenant boundaries that traditional scanners miss completely.

HOW WE SOLVE IT

Ground-Truth Matrix

Learns real resource ownership without guessing IDs, then compares cross-user responses using weighted Jaccard differential distance.

WHY WE WIN

Zero False Positives

Empirical live re-probes verify 100% reproducibility, generating shell-safe copy-paste cURL PoCs with zero cleartext token leaks.

Lead Security Architect: **Tanuj** | Full-Stack & Security Engineering

<https://github.com/x03tanuj/sentinel-api>

The Multi–Billion Dollar API Authorization Blindspot

⚠ Over 40% of all cloud API data breaches originate from **Broken Object Level Authorization (BOLA/IDOR)** — yet 99% of automated security scanners cannot detect it.



THE CORE FLAW

Changing `id=101` to `id=102`

An authenticated user simply changes an ID parameter in an API request and directly reads, modifies, or deletes another customer's private invoices, patient medical charts, or banking records.



TRADITIONAL TOOLS FAIL

DAST & SAST Are Blind to Auth

Conventional scanners only test for syntax injection like SQLi or XSS. Operating in single–user isolation, they have **zero understanding of multi–tenant ownership, role privileges, or cross–account access boundaries.**




SILENT EXPLOITATION

Flaws Evade Firewalls & WAFs

Because BOLA requests use valid JWT tokens and legitimate JSON formats, target APIs respond with standard **200 OK** payloads. WAFs see completely normal traffic, leaving breaches undetected for months.

The 4–Step Autonomous Differential Security Pipeline



Core Differentiator: Entirely autonomous from spec to PoC. No manual test writing, no synthetic dummy payloads, and zero secrets leaked.

Why SentinelAPI Wins: Standing Out Against 200 Teams

Judges will evaluate 200 teams on Track C. Most will submit simple prompt wrappers or basic injection fuzzers. Here is why SentinelAPI stands alone:

EVALUATION CRITERIA	TYPICAL HACKATHON TEAMS (200 TEAMS)	SENTINELAPI (OUR SOLUTION)
Vulnerability Focus	Generic SQLi, XSS, or syntax input fuzzing	Complex Multi-User Authorization (BOLA, BFLA, Data Exposure)
Detection Engine	Asks an LLM "is this vulnerable?" (Hallucinations & Flaky)	100% Deterministic Differential Engine (Weighted Jaccard Distance)
Ownership Knowledge	Random ID guessing or hardcoded IDs in test scripts	Autonomous Resource Ownership Discovery from Legitimate APIs
Accuracy & Evidence	Vague warnings with high false positive rates	Empirically Proven Live Reproduction + Copy-Paste cURL PoCs
Safety & Guardrails	Unrestricted flood requests; sends raw tokens to AI	Scope Guard, 20 RPS Rate Cap, Budget Cap, Zero Egress PII Guard
Production Readiness	Mock scripts or static Figma UI	FastAPI Backend + React 19 HUD + Headless CI Gate + Docker

The Core Security Engine: 6 Proprietary Deep-Tech Engines

ENGINE 1

Attack Surface Mapper

Parses OpenAPI 3.x specifications with circular `$ref` resolution. Extracts query/path parameters and scores initial endpoint risk based on state-changing operations.

ENGINE 2

Identity & HTTP Executor

Multi-persona JWT session manager enforcing 20 RPS token-bucket rate limits, 1,000 req budgets, and strict host allowlists (`assert_in_scope`).

ENGINE 3

Matrix Compiler

Discovers legitimate resource ownership per user and compiles a mathematical Allow/Deny expectation matrix for every endpoint and object ID.

ENGINE 4

Differential Analyzer

Computes weighted Jaccard payload similarity, HTTP status divergence, and detects leaked sensitive PII fields outside the documented schema.

ENGINE 5

6 Security Checks


Dedicated detection suites for BOLA/IDOR (API1), Broken Auth (API2), Data Exposure (API3), Rate Limiting (API4), BFLA Privileges (API5), and Input Anomalies (API8).

ENGINE 6

Explainable Risk Scorer

Replaces CVSS guesswork with a 4-part score: `Impact (40)` + `Exploitability (25)` + `Sensitivity (30)` + `Evidence (15)`.

Decoupled Architecture & Enforced Safety Boundaries



Core Safety Invariant: The deterministic security engine runs 100% offline. Zero credentials or tokens are ever persisted or exposed. The AI Analyst is strictly an evidence-only side-branch with zero raw data egress.

Shift-Left Automated Security Quality Gate for CI/CD

HEADLESS CI GATE CLI

Fails Builds on Violations

Integrates directly into GitHub Actions or GitLab CI. Breaks pull request builds with exit code 1 when CRITICAL or HIGH authorization vulnerabilities are detected.

SARIF CODE SCANNING

Native GitHub Security Alerts

Exports deterministic SARIF, Markdown, and JSON audit reports. Findings appear directly inside GitHub PR code review diffs with line-level remediation hints.

HERMETIC TEST SANDBOX

Verified Zero-Trust Baseline

Ships with realistic vulnerable sandboxes (Retail Store, Healthcare, FinTech). When **SECURE_MODE=true** is enabled, SentinelAPI verifies 100% controls pass.

GITHUB ACTIONS INTEGRATION (.github/workflows/ci.yml)

```
# 1. Trigger automated security scan against PR environment
python -m app.cli scan --spec http://api:8000/openapi.json --base-url http://api:8000 --json-out findings.json

# 2. Enforce quality gate (fails build on CRITICAL / HIGH findings with confidence ≥ 0.70)
python scripts/ci_gate.py --findings findings.json --fail-on CRITICAL,HIGH --min-confidence 0.70
```

Live Tactical Dashboard & Mathematical Differential Inspector

SentinelAPIv0.9.0-alpha

Scan History

Target: target_api:9000 (url:target_api:900...

Scope Guard: Active(allow-listed only)

SCANNING

Run Scan

WORKSPACE // ResultsTriage

target_api:9000

Findings

Authorization Matrix

Attack Surface

Analytics

Export

Explain top 5

AI calls: 0 / 15

AI Executive Summary

verify before use

Generate AI Summary

ENDPOINTS AUDITED

0 / 12100% surface mapped

CRITICAL FINDINGS

CLEAN0BOLA & write access

HIGH SEVERITY

CLEAN0State modification

MEDIUM SEVERITY

CLEAN0Data exposure

VERIFIED CONTROLS

27Expected 403 Denied

TARGET FINDINGS STREAM

0 of 0 Issues

All0Critical0High0Medium0

Low0Info0Secure

Search findings (e.g. /)

All Checks

Min Confidence:0%

Select a finding from the left rail to inspect differential evidence.

Results Triage Workspace: Real-Time Stream

Live telemetry HUD displaying verified control counts, severity badges, and one-click filtering across OWASP categories, status codes, and paths.

SentinelAPIv0.9.0-alpha

Scan History

Target: target_api:9000 (url:target_api:900...

Scope Guard: Active(allow-listed only)

READY

Run Scan

API5:2023by: userB

GET /users/{id}

Excessive Data Exposure on /users/{id}

API3:2023by: userA

GET /users/{id}

Excessive Data Exposure on /users/{id}

API3:2023by: userB

GET /orders/{id}

Broken Object Level Authorization (BOLA/IDOR) on /orders/{id}

API1:2023by: userB

GET /orders/{id}

Broken Object Level Authorization (BOLA/IDOR) on /orders/{id}

PII, credentials & financial records exposure

Empirical reproduction & differential confidence

bash - curl [Press 'c' to copy]

Copy cURL

Note: export TOKEN=<attacker token> first

curl -X GET http://target_api:9000/orders/101 -H "Authorization: Bearer \$TOKEN"

DUAL-IDENTITY RESPONSE DIFFERENTIAL

Left: Owner Baseline vs Right: Attacker Probe

Legitimate Owner (Legitimate Owner)200 OK

userB (Attacker Persona)Expected: 403 ForbiddenActual: 200 OK

{ "id": 101, "user_id": 1, "items": [{ "product_id": 1, "qty": 1 }, { "product_id": 2, "qty": 1 }], "total": 180, "shipping_address": "123 Elm St, Springfield, IL 62701", "status": "delivered" }

{ "id": 101, "user_id": 1, "items": [{ "product_id": 1, "qty": 1 }, { "product_id": 2, "qty": 1 }], "total": 180, "shipping_address": "123 Elm St, Springfield, IL 62701", "status": "delivered" }

"shipping_address": "12*****01" [LEAKED PII]

Differential Inspector: Side-by-Side Proof

Side-by-side baseline vs attack diff showing BOLA breach (200 OK vs 403), leaked PII fields, 4-factor risk breakdown, and copy-paste cURL PoC.

Empirical Evidence: Zero guesswork. Developers receive exact response diffs, 100% verified reproduction badges, and copy-paste shell commands.

Beyond the Hackathon: Path to an Enterprise SaaS Platform

PHASE A · COMMERCIAL SAAS

Open-Source CLI + Paid Cloud Dashboard

Maintain open-source CLI runner for individual developers and CI gates. Commercialize a hosted, multi-tenant SaaS dashboard for engineering organizations with centralized RBAC, historical risk trending, team collaboration, and vulnerability triage workflows.

PHASE B · ENTERPRISE AUTH

Enterprise OAuth2 & SSO Negotiation

Extend beyond static JWTs to automated OAuth2 / OIDC PKCE flow negotiation, SAML handshakes, and mTLS client-certificate authentication for complex enterprise banking, fintech, and healthcare environments.

PHASE C · COMPLIANCE MAPPINGS

Custom Regulatory Policy Engines

Provide declarative YAML policy mappings translating authorization findings directly into HIPAA, PCI-DSS, SOC 2, and ISO 27001 audit deliverables with automated executive PDF sign-off reports.

PHASE D · DISTRIBUTED MESH

Distributed Scanner Mesh & Persistence

Deploy distributed Kubernetes scanner workers capable of concurrently auditing large microservice meshes, backed by MongoDB and PostgreSQL persistence adapters for multi-year enterprise audit archiving.

Team & Project Summary — Questions & Live Demo

CORE TEAM & ARCHITECTURAL CONTRIBUTIONS

Tanuj — Lead Security Architect & Full-Stack Engineer

Amity University · AmiHacks 2026 (Track C: Cybersecurity)

- **Core Security Engine:** Differential response analyzer, ground-truth authorization matrix compiler, and 6 modular OWASP security checks.
- **Execution & Safety:** Token-bucket HTTP executor, strict immutable scope guard, rate/budget caps, and multi-persona JWT orchestrator.
- **Backend & Storage:** FastAPI async orchestrator, real-time SSE telemetry stream, and atomic JSON store with secret leak guards.
- **Tactical Web UI:** React 19/Vite dashboard, Results Triage Workspace, Authorization Heatmap, and Differential Inspector.
- **CI Quality Gate:** Headless `ci_gate.py`, Docker Compose testbed, live empirical reproduction, and zero-egress AI analyst boundary.

PROJECT SUMMARY & LIVE ACCESS

“SentinelAPI stops API authorization breaches before code hits production — fully autonomous, mathematically verified, with reproducible evidence and zero secrets leaked.”

GitHub Repository	https://github.com/x03tanuj/sentinel-api
Live Web Dashboard	http://localhost:8080
Automated Demo Script	<code>bash scripts/demo.sh</code>
Architecture Diagram	docs/architecture.png (2400×1400)

Thank You! We are now open for Q&A and Live Demo.